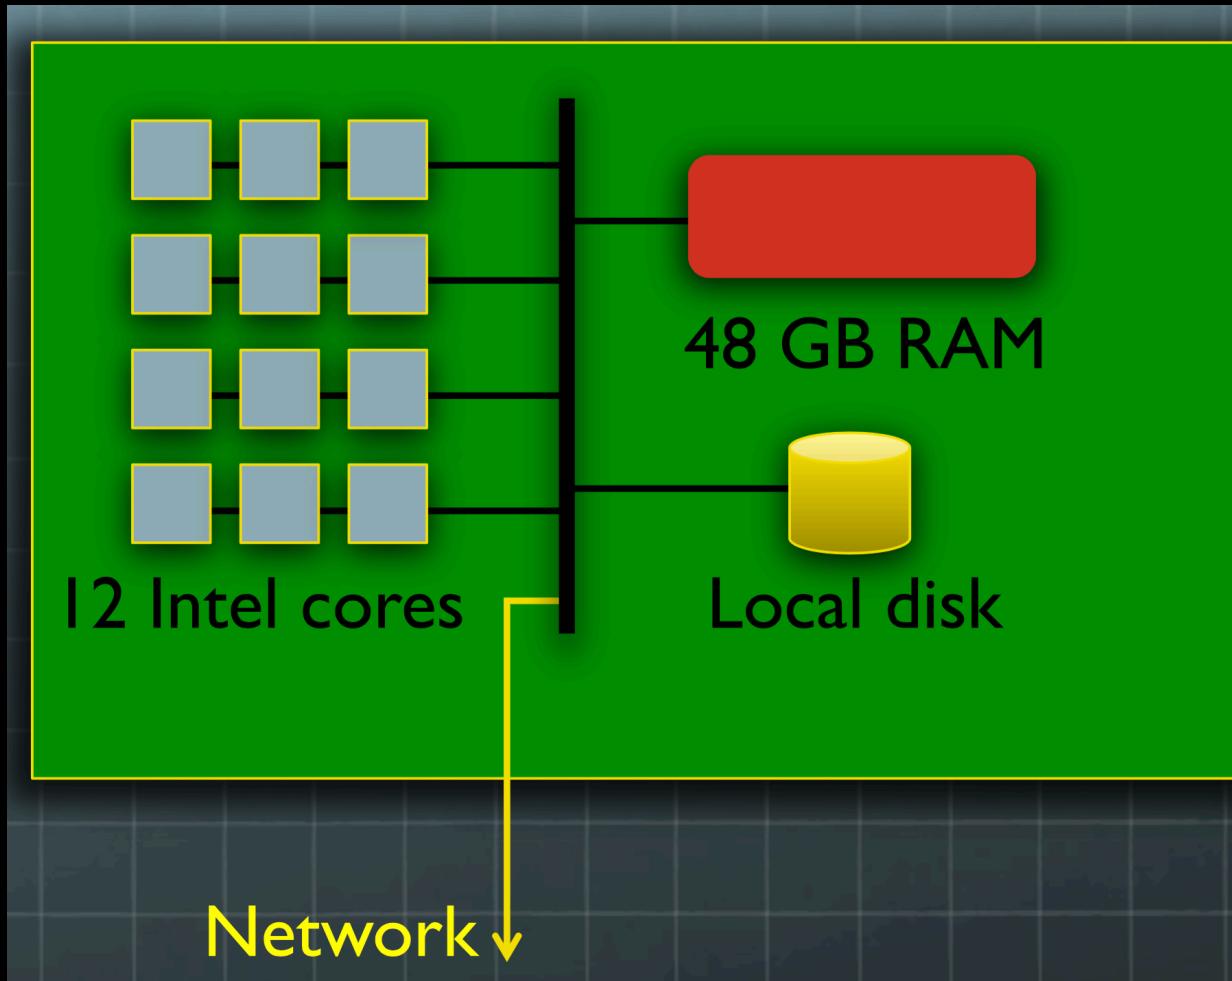
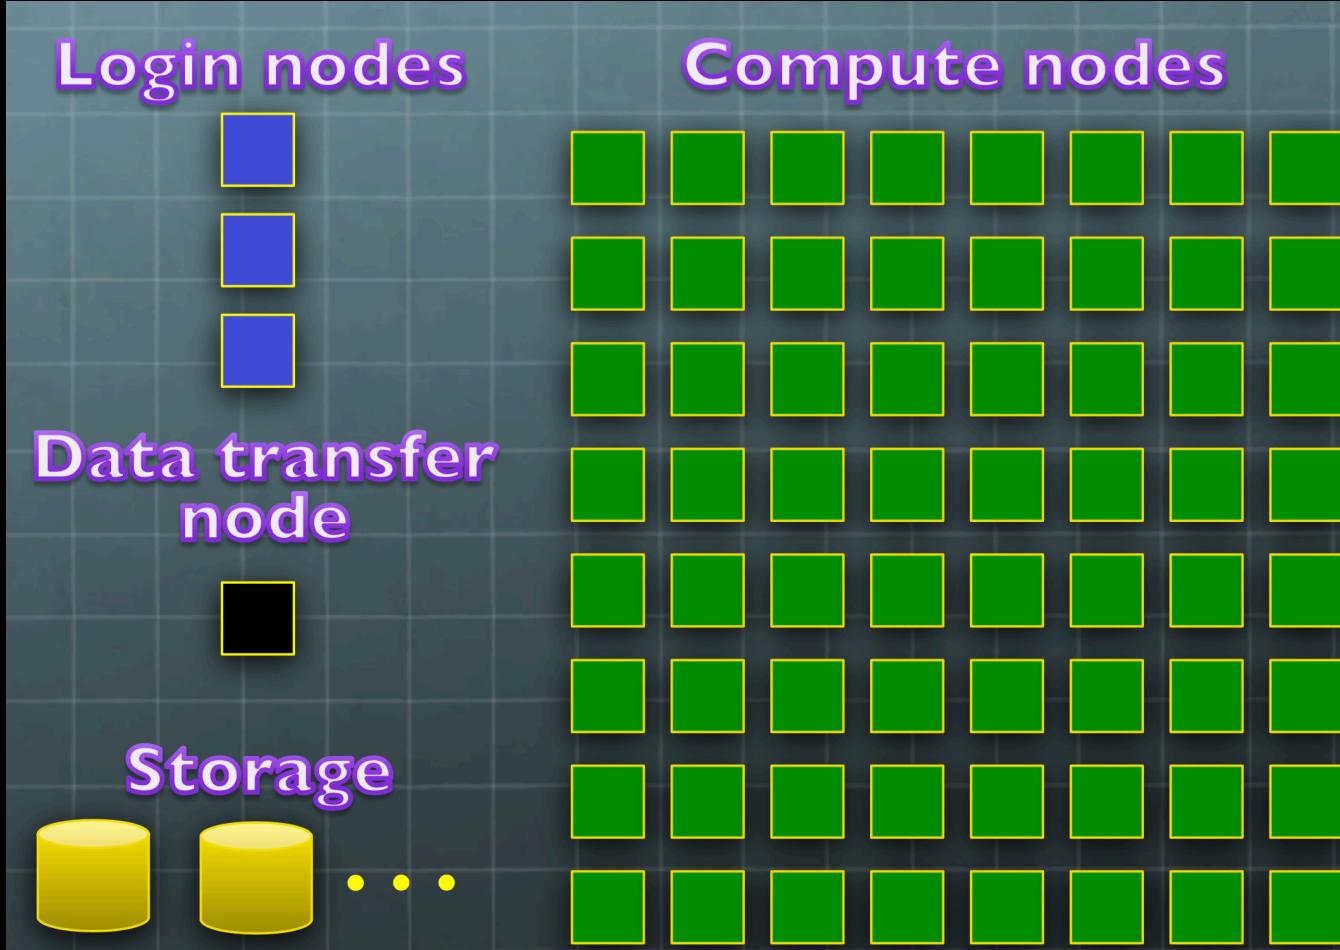


MODULE 03 – HIGH PERFORMANCE COMPUTING
HIGH-PERFORMANCE COMPUTING ON FLUX

A flux node



The flux cluster



To login the FLUX cluster...

- You need to be on-campus network
 - If you're off-campus, use VPN
<http://its.umich.edu/enterprise/wifi-networks/vpn>
- You need to enroll to Duo two-factor authorization
<http://its.umich.edu/accounts-access/uniqnames-passwords/two-factor-authentication>
- Follow these steps:
 1. Login to **[uniqname]@flux-login.engin.umich.edu** via SSH
 2. Type your Kerberos password
 3. Complete login through duo two-factor authorization

Writing PBS batch script

- Start from the guide

<http://arc-ts.umich.edu/flux-user-guide/>

<http://www.youtube.com/watch?v=SW8Lu1-JaSM>

<http://arc-ts.umich.edu/software/torque/pbs-template/>

- At the minimum you will need

PBS -V (to inherit the environment variables)

PBS -A bios815_flux (account info)

PBS -l qos=flux (quality of service, use flux for standard jobs)

PBS -q flux (queue of the jobs, use flux for standard jobs)

PBS -l nodes=1,pmem=4gb,walltime=4:00:00

(specify the resources you need)

Running interactive jobs

```
[hmkang@flux-login3 815]$ cat minimum.pbs.sh
#!/bin/bash
#PBS -V
#PBS -A sph_flux
#PBS -l qos=flux
#PBS -q flux
#PBS -l nodes=1,pmem=1gb,walltime=4:00:00
[hmkang@flux-login3 815]$ qsub -I minimum.pbs.sh
qsub: waiting for job 21743003.nyx.arc-ts.umich.edu to start
qsub: job 21743003.nyx.arc-ts.umich.edu ready

[hmkang@nyx6222 ~]$
```

Scratch file system

- **/home directory in FLUX has limited amount of space available.**
 - Store important scripts, configurations that need to be backed up
- Any large data files must be stored under
/scratch/[pi_uniqname]_flux/[your_uniqname]/
 - These files will not be backed up, and transfer the final outcome to a safe storage

```
[hmkang@flux-login3 815]$ df -h /home
```

Filesystem	Size	Used	Avail	Use%	Mounted on
nyxhome.value.storage.umich.edu:/nyxhome	23T	22T	1.3T	95%	/home

```
[hmkang@flux-login3 815]$ df -h /scratch
```

Filesystem	Size	Used	Avail	Use%	Mounted on
10.255.7.0@o2ib:10.255.7.1@o2ib:/scratch	1.5P	1.3P	201T	87%	/scratch

Example 1 : Gaussian Mixture LRT

- Consider the following hypothesis testing procedure

$$H_0 : x_i \sim N(\mu, \sigma^2)$$

$$H_1 : x_i \sim \phi N(\mu_1, \sigma_1^2) + (1 - \phi) N(\mu_2, \sigma_2^2)$$

where x_1, x_2, \dots, x_n are i.i.d.

In other words, we are testing whether $\pi = 0$ or not.

- Consider a likelihood-ratio test between these two hypotheses.
- We want to calculate the power of the LRT at size α at a particular configuration of $\theta = (\phi, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2)$

Constructing a LRT

$$L_0 = L(\hat{\theta}_0 | \mathbf{x}) = \prod_{i=1}^n \left[\frac{1}{\sqrt{2\pi\hat{\sigma}^2}} \exp\left(-\frac{(x_i - \hat{\mu})^2}{2\hat{\sigma}^2}\right) \right]$$

$$\hat{\mu} = \bar{x}, \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

$$L_1 = L(\hat{\theta} | \mathbf{x}) = \max_{(\phi, \mu_1, \mu_2, \sigma_1^2, \sigma_2^2)} \prod_{i=1}^n \left[\frac{\phi}{\sqrt{2\pi\sigma_1^2}} \exp\left(-\frac{(x_i - \mu_1)^2}{2\sigma_1^2}\right) + \frac{1 - \phi}{\sqrt{2\pi\sigma_2^2}} \exp\left(-\frac{(x_i - \mu_2)^2}{2\sigma_2^2}\right) \right]$$

$$T = 2[\log(L_1) - \log(L_0)] \sim ?$$

Three key questions

- **How should we compute T and L_1 ?**
 - Calculating L_1 is not straightforward analytically
- **How can we construct a size α test?**
 - The asymptotic distribution of T is not straightforward to obtain
- **How should we compute the power?**
 - The closed form of the power function is difficult to obtain

Answers are “numerical methods”

- **How should we compute T and L_1 ?**
 - Use E-M algorithm!
- **How can we construct a size α test?**
 - Obtain the empirical null distribution via simulation
- **How should we compute the power?**
 - Use simulation to compute the power

Implementing power simulation in R

```
em2LRT <- function(x) {  
  w <- rbinom(n.sample, 1, 0.5) ## initial assignment  
  means <- sample(x,2)  
  sds <- rep(sd(x),2)  
  for(i in 1:200) {  
    llks <- cbind( dnorm(x, means[1], sds[1]) * phi,  
                  dnorm(x, means[2], sds[2]) * (1-phi) ) + 1e-100  
    w <- llks[,1] / rowSums(llks)  
    phi <- mean(w)  
    means <- c( mean(w * x)/phi, mean( (1-w)*x ) / (1-phi) )  
    sds <- sqrt( c( mean( w * ( x - means[1] )^2 ) / phi,  
                  mean( (1-w)* ( x - means[2] )^2 ) / (1-phi) ) )  
  }  
  llk.alt <- sum(log(rowSums(cbind( dnorm(x, means[1], sds[1])*phi,  
                                dnorm(x, means[2], sds[2])*(1-phi) ) ) + 1e-100))  
  llk.null <- sum(log(dnorm(x, mean(x), sqrt(mean(x^2) - mean(x)^2))+1e-100))  
  llk.diff <- llk.alt - llk.null  
  if ( llk.diff < 0 ) { llk.diff <- 0 }  
  return(list(llk1=llk.alt, llk0=llk.null, stat=2*llk.diff,  
             log10p = pchisq(2*llk.diff, df=1, lower.tail=FALSE, log.p=TRUE)/log(10)))  
}
```

Handling arguments and running simulations

```
args <- commandArgs(TRUE)
out <- args[1]
phi <- as.double(args[2])
mul <- as.double(args[3])
mu2 <- as.double(args[4])
sd1 <- as.double(args[5])
sd2 <- as.double(args[6])
n.sample <- as.integer(args[7])
n.simul <- as.integer(args[8])
rst <- matrix(NA, n.simul, 4)
for(i in 1:n.simul) { ## perform simulation and record statistics
  w <- rbinom(n.sample,1,phi)
  x <- rnorm(n.sample, mul, sd1) * w + rnorm(n.sample, mu2, sd2) * (1-w)
  r <- em2LRT(x)
  rst[i,] <- c(r$llk1,r$llk0,r$stat,r$log10p)
}
colnames(rst) <- c("LLK1","LLK0","LRT","LOG10P");
write.table(rst,out,row.names=FALSE,col.names=TRUE,quote=FALSE,sep="\t");
```

Running the simulation

```
$ /usr/bin/time -v Rscript gmmpower.r tmp.out 0.5 0 0 1 1 1000 1000
Command being timed: "Rscript gmmpower.r tmp.out 0.5 0 0 1 1 1000 1000"
User time (seconds): 54.31
System time (seconds): 0.04
Percent of CPU this job got: 99%
Elapsed (wall clock) time (h:mm:ss or m:ss): 0:54.40
Average shared text size (kbytes): 0
Average unshared data size (kbytes): 0
Average stack size (kbytes): 0
Average total size (kbytes): 0
Maximum resident set size (kbytes): 25840
...
...
```

Power calculation via simulation

- To calculate the power, we need to estimate

$$\Pr(T > t_\alpha | \theta)$$

where

$$\Pr(T > t_\alpha | \theta_0) = \alpha$$

- Suppose that we have **n** samples from the null distribution (**X**) and **m** samples from the alternative distribution (**Y**). We estimate the power using

$$\hat{\beta} = \frac{\sum_{i=1}^n I(Y_i > X_{(n-n\alpha+1)})}{m}$$

Example power calculation

- For obtaining the null distribution..
 - $n = 1,000,000$
- For each alternative hypothesis...
 - $m = 10,000$
- How many E-M iterations are required?
- How long would the job take if it runs with a single core?

Running one batch of 10,000 simulations

```
#!/bin/bash

#PBS -V
#PBS -A hmkang_flux
#PBS -l qos=flux
#PBS -q flux
#PBS -l nodes=1,pmem=1gb,walltime=4:00:00
export WD=/scratch/hmkang_flux/hmkang/class/815
Rscript ${WD}/gmmpower.r ${WD}/out1.txt 0.5 0 0 1 1
1000 10000
```

Use **qsub foo.pbs.sh** to submit the job to the cluster

Ways to run 1,000,000 simulations

(in 10 minutes)

- 1. Create 100 different PBS scripts,
with different output file names,
and submit each PBS script using qsub**

- 2. Create one PBS script that can take an argument, and run the
script 100 times with different arguments**

- 3. Create a job array script and run the script once, specifying the
number of simulations to run**

Creating a job array script

```
#!/bin/bash
#PBS -V
#PBS -A hmkang_flux
#PBS -l qos=flux
#PBS -q flux
#PBS -t 1-100
#PBS -l nodes=1,pmem=1gb,walltime=4:00:00
export WD=/scratch/hmkang_flux/hmkang/class/815
Rscript ${WD}/gmmpower.r
${WD}/array.${PBS_ARRAYID}.out 0.5 0 0 1 1 1000
10000
```

See <http://arc-ts.umich.edu/software/torque/job-arrays/> for more details

Submitting the array job

- Simply run

```
$ qsub gmmpower.array.sh
```

- Check the status of the job using

```
$ qstat | grep [your_uniqname] , or  
$ qstat -f [job_id]
```

Passing arguments by environment variables

```
#!/bin/bash
#PBS -V
#PBS -A hmkang_flux
#PBS -l qos=flux
#PBS -q flux
#PBS -l nodes=1,pmem=1gb,walltime=4:00:00
export
WD=/scratch/hmkang_flux/[your_uniqname]/class/815
Rscript ${WD}/gmmpower.r ${WD}/${out} 0.5 0 ${mu} 1 1
1000 10000
```

Submitting 100 jobs with different arguments

```
$ seq -w 1 100 | xargs -I {} qsub -v  
out='null.{}.out',mu='0' gmpower.pbs.sh
```

- What do you think it will happen?
- For running jobs under the alternative hypothesis, what do we need to do?

Running simulations under the alternative hypotheses

```
$ seq -w 0 0.1 2 | xargs -I {} qsub -v  
out='alt.{}.out',mu='{}' gmpower.pbs.sh
```

- What do you think it will happen?
- Can you do similar tasks using job array?

Merging the null distribution

- Merge multiple files can be done with simple UNIX command in many cases

```
$ cat null.*.out | grep -v ^#LLK1 >  
merged.null.out
```

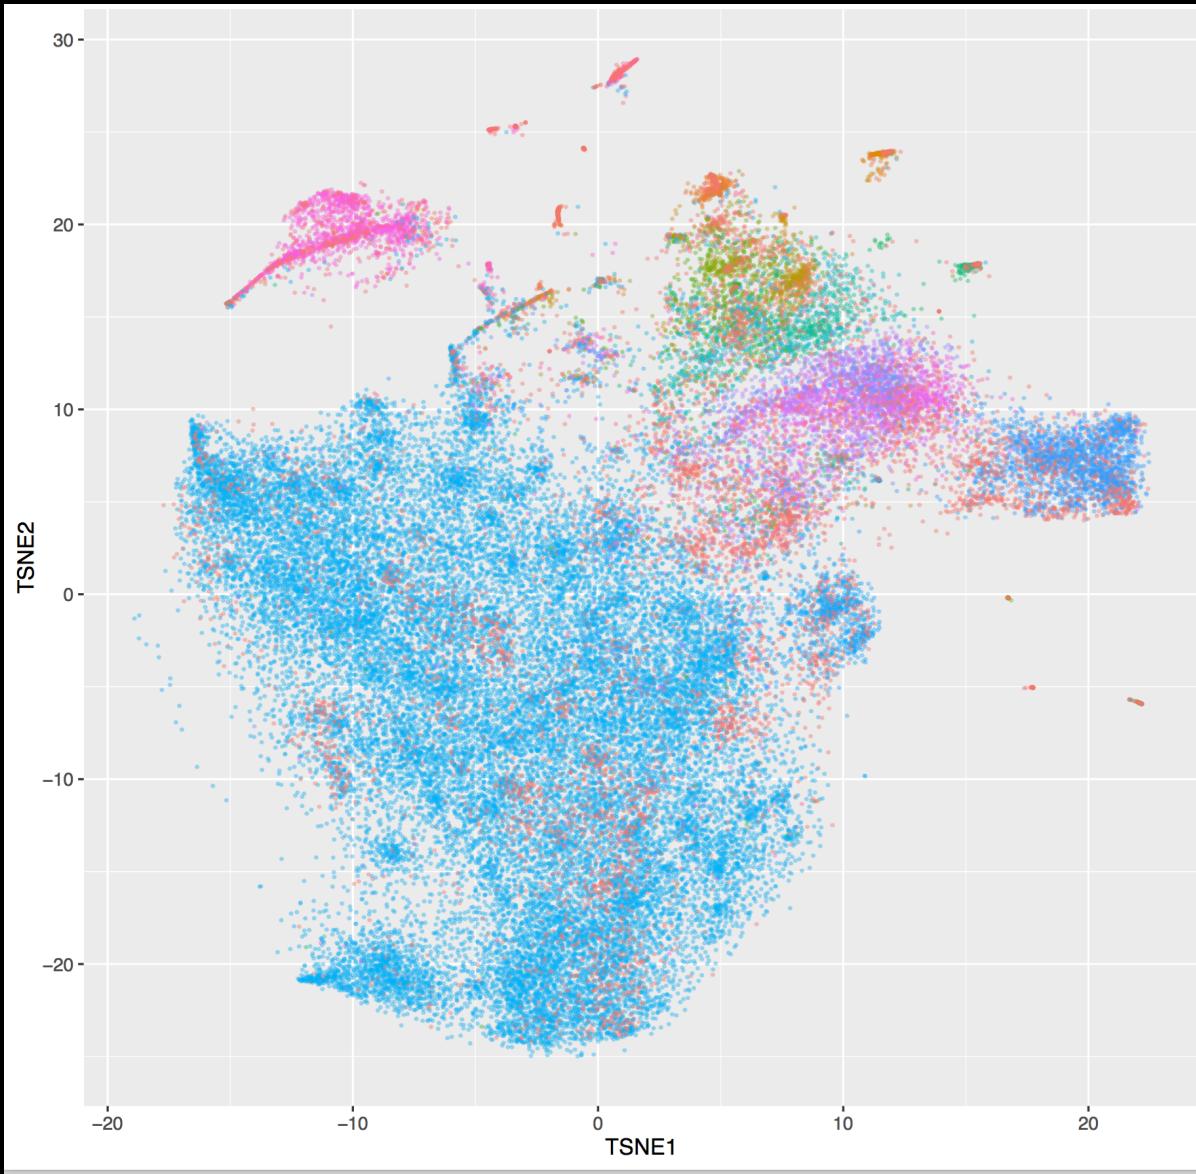
- Automating the following analysis is also possible, but it's rather trivial tasks.

Handling Big Data

```
/scratch/hmkang_flux/hmkang/class/815/GSM1626793_16  
26799_P14Retina_merged.digital.expression.txt.gz  
/scratch/hmkang_flux/hmkang/class/815/GSM1626793_16  
26799_P14Retina_merged.digital.expression.tsne.dat
```

These files contains number of RNA transcript in mouse retina measured for each of 49,300 cell and 24,760 transcripts (Macosko et al. 2015)

File size is only 52M (sparse), but contains matrix with 1.2 billion elements



What we want is

- Calculate the p-value for each gene in terms of association with TSNE1 and TSNE2 coordinates, accounting for the difference of total read counts for each cell.
- ...Using either negative binomial, or Poisson regression.

Reviewing what we have learned..

- Numerical Methods
- Algorithms and data structures
- High-performance computing